

Extra Credit Opportunity

Prof. Tianbao Yang, University of Iowa

Lecture 12: Artificial Neural Networks (Mitchell Chapter 4)

CS 167: Machine Learning

Deep Learning with Big and Small Data

Deep learning has brought tremendous success in many areas with the help of big data and super computing. In this talk, I will present the state of art results of deep learning for image classification. I will also talk about our recent research on how to learn a deep convolutional neural network for fine-grained image classification where big labeled data is difficult to be obtained.

Friday, October 21, in Meredith 106 at 2:00pm.

Earn 3 extra credit (homework) points for


- attend the talk
- write up a paragraph on something you learned

CS 167: Machine Learning

L12: ANN

2 / 15

Reminder: Perceptron Unit



$$o(x_1, \dots, x_n) = \begin{cases} 1 & \text{if } w_0 + w_1 x_1 + \dots + w_n x_n > 0 \\ -1 & \text{otherwise.} \end{cases}$$

Sometimes we'll use simpler vector notation:



$$o(\vec{x}) = \begin{cases} 1 & \text{if } \vec{w} \cdot \vec{x} > 0 \\ -1 & \text{otherwise.} \end{cases}$$

CS 167: Machine Learning

L12: ANN

3 / 15

Multilayer Networks




input layer → hidden layer → output layer

CS 167: Machine Learning


L12: ANN

4 / 15


Expressiveness of Multilayer Networks



Single perceptron decision surface



2 opposite thresholds: ridge




2 opposite ridges: bump

Multilayer Perceptron Classification Example


Classify the following four examples (on two attributes A_1 and A_2) using the given multilayer perceptron artificial neural network.

example #	A_1	A_2
1	3	5
2	2	7
3	1	1
4	2	3



Exercise: Do examples 2-4.

Sigmoid Unit



$\sigma(x)$ is the sigmoid function

$$\frac{1}{1 + e^{-x}}$$

Nice property: $\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$

We can derive gradient decent rules to train

- One sigmoid unit
- *Multilayer networks* of sigmoid units → Backpropagation
- Read textbook for derivation of training rule

Backpropagation Algorithm

Initialize all weights to small random numbers.

Until satisfied, Do

- For each training example, Do
 - 1 Input the training example to the network and compute the network outputs (o_u for unit u)
 - 2 For each output unit k

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$

- 3 For each hidden unit h

$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{h,k} \delta_k$$

- 4 Update each network weight $w_{i,j}$

$$w_{i,j} \leftarrow w_{i,j} + \Delta w_{i,j}$$

where

$$\Delta w_{i,j} = \eta \delta_j x_{i,j}$$

Backpropagation of the error

The error function in the Perceptron/SGD training rule was $(t_k - o_k)$:

$$w_i \leftarrow w_i + \Delta w_i$$
$$\Delta w_i = \eta(t - o)x_i$$

The Backpropagation training rule is more complex: $o_k(1 - o_k)(t_k - o_k)$

Why? $o_k(1 - o_k)$ is the derivative of the sigmoid squashing function, important for deriving the training rule.

Both rules scale based on the learning rate, η and the input value x_{ij}

For hidden units, the error at the output unit is divided up among the units feeding into it based on their weight - how much each is responsible for the error at the output unit

Discussion Questions

How many times do you do each training example?

What should η be? Can/should it change during training?

How long does training take?

How long does it take to classify a new example?

Can we get unlucky with initial random weights?

Can we do regression?

Expressive Capabilities of ANNs

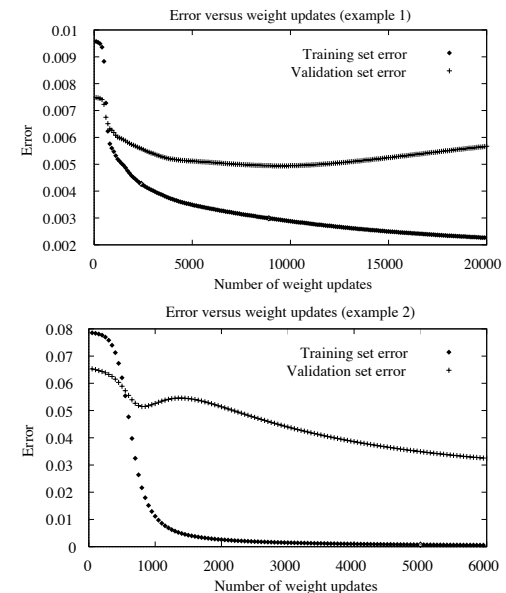
Boolean functions:

- Every boolean function can be represented by network with single hidden layer
- but might require exponential (in number of inputs) hidden units

Continuous functions:

- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers

Overfitting in ANNs




Recurrent Networks


Let's update your version of scikit-learn

It seems that there's a lot of good stuff in scikit-learn 0.18, so let's update it if you haven't already:


- 1 In Anaconda Navigator, select [Environments](#) on the left-hand side menu
- 2 Click the [Update Index](#) button across the top of the window
- 3 Scroll to find scikit-learn in the list of packages (or search for it)
- 4 Look under the version column, hopefully it's blue with an up-right arrow. That means an upgrade is available.
- 5 If so, right/crtl-click the package and select [mark for update](#).
- 6 Click the [Apply](#) button near the bottom right of the window.



(a) Feedforward network



(b) Recurrent network



(c) Recurrent network unfolded in time

Exercise

Determine the answer to these questions:

- Does [scikit-learn](#) have a neural network for either [classification](#) or [regression](#)?
- If so, what is the default network structure it uses (how many layers?, how many nodes in each layer?)? How can I change it?
- If so, does it use the sigmoid squashing (activation) function? If not, can I set it to use it?
- If so, what is the default learning rate and number of epochs? How can you change them?

Exercises:

- What's the smallest network that performs well on the Iris data?
- Can you make an ANN that beats the SGDRegressor performance on the MPG data?