

Lecture 13: Face Recognition Lab

CS 167: Machine Learning

Face Recognition

Face Recognition: classification task - identify a person from a photograph of their face

Our Data Set: Labeled Faces in the Wild from University of Massachusetts

<http://vis-www.cs.umass.edu/lfw/>



CS 167: Machine Learning

L13: Face Recognition Lab

2 / 8

Labeled Faces in the Wild

All photos in this set are 250×250 pixels

Each pixel is an attribute: lots of attributes 250×250 pixels * 3 colors = 1875000

Convert to grayscale: 62500 total attributes

scikit-learn will load this dataset for you: by default converts to grayscale and crops/resizes to 62×47

Let's do it

```
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.cross_validation import train_test_split
from sklearn.datasets import fetch_lfw_people
from sklearn.metrics import accuracy_score
from sklearn.neural_network import MLPClassifier
import numpy

lfw_people = fetch_lfw_people(min_faces_per_person=70)
print(len(lfw_people.images))
print(lfw_people.target_names)
```

How many people do they have 40 or more photos of?

Looking at some of the data

```
image_num = 0 #try plugging in some other numbers too

print("image: ",lfw_people.images[image_num])

plt.imshow(lfw_people.images[image_num],cmap='gray')
plt.show()

img_target = lfw_people.target[image_num]
print(img_target)
print(lfw_people.target_names[img_target])

print("data: ", lfw_people.data[image_num])
print("data length: ",len(lfw_people.data[image_num]))
```

What's the difference between

`lfw_people.images` and `lfw_people.data`?

`lfw_people.target` and `lfw_people.target_names`?

Training, Testing

Note: the `train_test_split` is a little different because the predictive columns and target column are not in the same dataframe - actually, it's not even a dataframe but a `numpy` array

```
(train_data, test_data, train_target, test_target)\
    = train_test_split(lfw_people.data, lfw_people.target,\
                        test_size = 0.2)
mlp_classifier = MLPClassifier()
mlp_classifier.fit(train_data,train_target)
predictions = mlp_classifier.predict(test_data)
print(accuracy_score(predictions,test_target))
```

How did we do? Is that good? Is it overfitting?

See next slide for code to create a gallery of your predictions

Prediction Gallery

```
for i in range(len(predictions)):
    plt.imshow(test_data[i].reshape(62,47),cmap='gray')
    plt.show()
    print("predicted: ",\
          lfw_people.target_names[predictions[i]],\
          " actual: ", lfw_people.target_names[test_target[i]])
```

since we didn't split the `lfw_people.images` version, we have to reshape the `data` array to be like the `image` 2D array

Exercise: Tweak the Neural Network parameters until you start getting good predictions. Make sure to try different network structures (number of layers and nodes within each layer)

Let's figure out some other data

I've heard that these data sets don't necessarily work for linear models:

<https://archive.ics.uci.edu/ml/datasets/Contraceptive+Method+Choice>

<https://archive.ics.uci.edu/ml/datasets/Forest+Fires>

Try the following

- Get a baseline accuracy/MSE with a linear model
- Try to find a good neural net.